

```
In [329]: import matplotlib.pyplot as plt
import numpy as np
import scipy as sci
import pandas as pd
import seaborn as sns
import matplotlib
```


```
In [330]: train = pd.read_csv("house-prices-advanced-regression-techniques/train.csv")
test = pd.read_csv("house-prices-advanced-regression-techniques/test.csv")
```

```
In [331]: train.head()
```

Out[331]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Util
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	Al
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	Al
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	Al
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	Al
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	Al

5 rows × 81 columns



```
In [332]: data = pd.concat((train.loc[:, 'MSSubClass': 'SaleCondition'], test.loc[:, 'MSSubClass': 'SaleCondition']))
```

```
In [333]: train["SalePrice"] = np.log1p(train["SalePrice"])

#Log transform skewed numeric features:
indeces = data.dtypes[data.dtypes != "object"].index

skewed_feats = train[indeces].apply(lambda x: skew(x.dropna())) #compute skewness
skewed_feats = skewed_feats[skewed_feats > 0.75]
skewed_feats = skewed_feats.index

all_data[skewed_feats] = np.log1p(all_data[skewed_feats])
```

```
In [334]: data = pd.get_dummies(data)
```

```
In [335]: #filling NA's with the mean of the column:
data = data.fillna(data.mean())
```

```
In [336]: #creating matrices for sklearn:
X_train = all_data[:train.shape[0]]
X_test = all_data[train.shape[0]:]
Y_train = train.SalePrice
```

```
In [337]: from sklearn.linear_model import Ridge, Lasso
          from sklearn.model_selection import cross_val_score
```

```
In [338]: model_ridge = Ridge(alpha=.1)
          model_ridge.fit(X_train, Y_train)
```

Out[338]: Ridge(alpha=0.1)

```
In [339]: preds = model_ridge.predict(X_test)
          preds = np.expm1(preds)
```

```
In [340]: prediction = pd.DataFrame({"id":test.Id, "SalePrice":preds})
          prediction.to_csv("ridge_sol.csv", index = False)
```


In [341]: *#After submitting to Kaggle, we get a RMSE of .1377*

```
In [342]: def rmse_cv(model):
          rmse= np.sqrt(-cross_val_score(model, X_train, Y_train, scoring="neg_mean_
          squared_error", cv = 5))
          return(rmse)
```

```
In [343]: alphas = [0.05, 0.1, 0.3, 1, 3, 5, 10, 20, 30]
          cv_ridge = [rmse_cv(Ridge(alpha = alpha)).mean()
                      for alpha in alphas]
```

```
In [344]: cv_ridge = pd.Series(cv_ridge, index = alphas)
          cv_ridge.plot(title = "Validation Ridge")
          plt.xlabel("alpha")
          plt.ylabel("rmse")
```

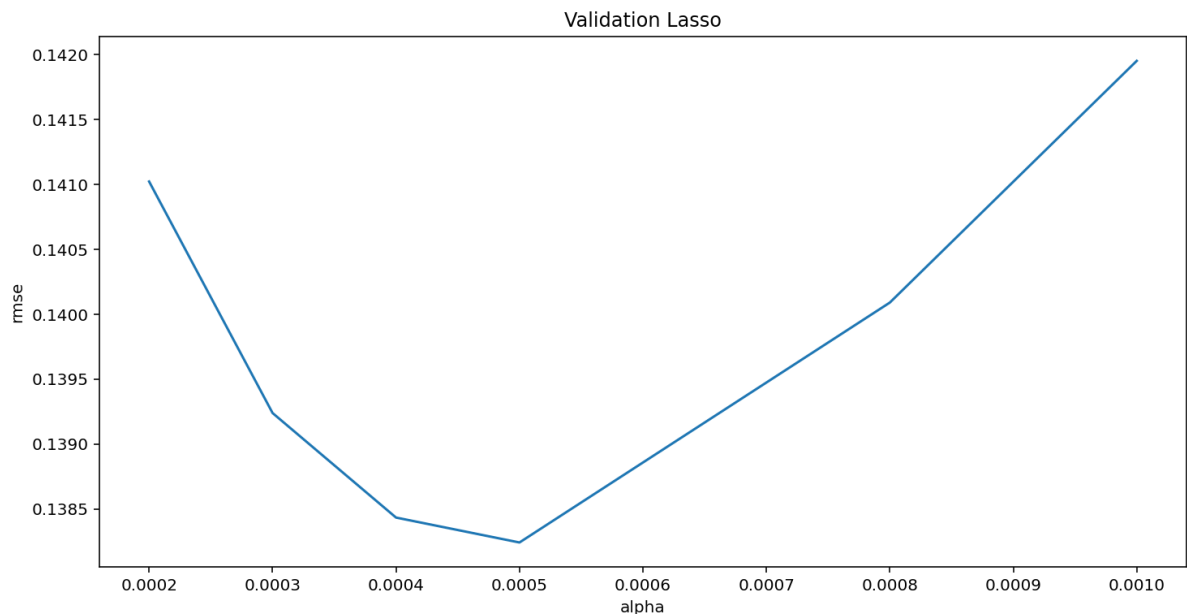
Out[344]: Text(0, 0.5, 'rmse')



```
In [345]: alphas1 = [.001, .0008, .0005, .0004, .0003, .0002]
cv_lasso = [rmse_cv(Lasso(alpha = alpha)).mean()
             for alpha in alphas1]
```

```
In [346]: cv_lasso = pd.Series(cv_lasso, index = alphas1)
cv_lasso.plot(title = "Validation Lasso")
plt.xlabel("alpha")
plt.ylabel("rmse")
```

Out[346]: Text(0, 0.5, 'rmse')



```
In [347]: # For a single LASSO Model, we can get to a RMSE of ~.138
# For a single Ridge Model, we can get to a RMSE of ~.141
```

```
In [348]: models_lasso = [Lasso(alpha = alpha).fit(X_train, Y_train) for alpha in alphas1]
```

```
In [349]: coefs = [pd.Series(models_lasso[i].coef_, index = X_train.columns) for i in range(0, len(alphas1))]
```


```
In [350]: l0 = np.zeros_like(alphas1)
for i in range(0, len(alphas1)):
    l0[i] = sum(coefs[i] != 0)
```

```
In [351]: print(l0)
```

```
[ 80.  90. 109. 119. 132. 157.]
```

```
In [352]: l0 = pd.Series(l0, index = alphas1)
l0.plot(title = "Norm")
plt.xlabel("alpha")
plt.ylabel("number of coefs")
```

```
Out[352]: Text(0, 0.5, 'number of coefs')
```



```
In [353]: predictions_lasso = [models_lasso[i].predict(X_train) for i in range(0, len(alphas1))]
```

```
In [358]: X_train.loc[:, "lasso1"] = predictions_lasso[0]
X_train.loc[:, "lasso2"] = predictions_lasso[1]
X_train.loc[:, "lasso3"] = predictions_lasso[2]
X_train.loc[:, "lasso4"] = predictions_lasso[3]
X_train.loc[:, "lasso5"] = predictions_lasso[4]
X_train.loc[:, "lasso6"] = predictions_lasso[5]
```

```
In [363]: model_ridge_es = Ridge(alpha=10)
model_ridge_es.fit(X_train, Y_train)
```

```
Out[363]: Ridge(alpha=10)
```

```
In [364]: def rmse_cv(model):
    rmse= np.sqrt(-cross_val_score(model, X_train, Y_train, scoring="neg_mean_squared_error", cv = 5))
    return(rmse)
```

```
In [366]: cv_ridge_es = rmse_cv(Ridge(alpha = 10)).mean()
print(cv_ridge_es)
```

```
0.13445149400996287
```

```
In [ ]: # We can get down to a RMSE score of .134 which is better than both the LASSO
        and the Ridge Models
```