

```
In [2]: import torch
import torchvision
import torchvision.transforms as transforms
import numpy as np
```

```
In [3]: transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = torchvision.datasets.CIFAR10(root='/home/CIFAR-10 Classifier Using
    CNN in PyTorch/data/',
                                       train=True,
                                       download=True,
                                       transform=transform)
trainloader = torch.utils.data.DataLoader(trainset,
                                           batch_size=4,
                                           shuffle=True)
testset = torchvision.datasets.CIFAR10(root='./data',
                                       train=False,
                                       download=True,
                                       transform=transform)
testloader = torch.utils.data.DataLoader(testset,
                                          batch_size=4,
                                          shuffle=False)

classes = ('plane', 'car', 'bird', 'cat', 'deer',
           'dog', 'frog', 'horse', 'ship', 'truck')
```

Files already downloaded and verified

Files already downloaded and verified

```
In [17]: def convert_to_imshow_format(image):
# first convert back to [0,1] range from [-1,1] range
image = image / 2 + 0.5
image = image.numpy()
# convert from CHW to HWC
# from 3x32x32 to 32x32x3
return image.transpose(1,2,0)

dataiter = iter(trainloader)
images, labels = dataiter.next()

fig, axes = plt.subplots(1, len(images), figsize=(12,2.5))
for idx, image in enumerate(images):
    axes[idx].imshow(convert_to_imshow_format(image))
    axes[idx].set_title(classes[labels[idx]])
    axes[idx].set_xticks([])
    axes[idx].set_yticks([])
```



```
In [4]: import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
```

In [5]: `import torch.optim as optim`

```
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

In [6]: `import os`

```
model_directory_path = '/home/CIFAR-10 Classifier Using CNN in PyTorch/model/'
model_path = model_directory_path + 'cifar-10-cnn-model.pt'

if not os.path.exists(model_directory_path):
    os.makedirs(model_directory_path)

if os.path.isfile(model_path):
    # Load trained model parameters from disk
    net.load_state_dict(torch.load(model_path))
    print('Loaded model parameters from disk.')
else:
    for epoch in range(2): # Loop over the dataset multiple times

        running_loss = 0.0
        for i, data in enumerate(trainloader, 0):
            # get the inputs
            inputs, labels = data

            # zero the parameter gradients
            optimizer.zero_grad()

            # forward + backward + optimize
            outputs = net(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            # print statistics
            running_loss += loss.item()
            if i % 2000 == 1999: # print every 2000 mini-batches
                print('[%d, %5d] loss: %.3f' %
                      (epoch + 1, i + 1, running_loss / 2000))
            running_loss = 0.0
        print('Finished Training.')
        torch.save(net.state_dict(), model_path)
        print('Saved model parameters to disk.')
```

Loaded model parameters from disk.

```
In [18]: dataiter = iter(testloader)
         images, labels = dataiter.next()

         fig, axes = plt.subplots(1, len(images), figsize=(12,2.5))
         for idx, image in enumerate(images):
             axes[idx].imshow(convert_to_imshow_format(image))
             axes[idx].set_title(classes[labels[idx]])
             axes[idx].set_xticks([])
             axes[idx].set_yticks([])
```



```
In [8]: outputs = net(images)
```

```
In [9]: outputs
```

```
Out[9]: tensor([[ -0.5587, -1.3294, -0.0147,  1.7789, -0.9728,  0.8734,  0.2706, -0.60
      86,
               -0.2059,  0.2847],
               [ 4.8533,  5.2669, -0.6852, -2.5590, -2.6156, -4.4962, -2.6582, -4.18
      07,
               5.4332,  3.9467],
               [ 2.4263,  2.5022, -0.4277, -1.3331, -1.3506, -2.5808, -1.3142, -2.63
      37,
               3.4305,  2.0966],
               [ 3.4028,  1.2159,  0.6023, -1.1971, -1.3105, -2.9989, -1.2615, -2.89
      25,
               4.5330,  0.9981]], grad_fn=<AddmmBackward>)
```

```
In [10]: sm = nn.Softmax(dim=1)
         sm_outputs = sm(outputs)
         print(sm_outputs)

         tensor([[3.9401e-02, 1.8229e-02, 6.7883e-02, 4.0804e-01, 2.6041e-02, 1.6498e-
      01,
               9.0294e-02, 3.7484e-02, 5.6070e-02, 9.1576e-02],
               [2.1242e-01, 3.2121e-01, 8.3527e-04, 1.2825e-04, 1.2120e-04, 1.8481e-
      05,
               1.1614e-04, 2.5336e-05, 3.7934e-01, 8.5787e-02],
               [1.7642e-01, 1.9033e-01, 1.0164e-02, 4.1103e-03, 4.0391e-03, 1.1803e-
      03,
               4.1887e-03, 1.1195e-03, 4.8157e-01, 1.2688e-01],
               [2.2771e-01, 2.5563e-02, 1.3840e-02, 2.2892e-03, 2.0437e-03, 3.7772e-
      04,
               2.1464e-03, 4.2012e-04, 7.0505e-01, 2.0561e-02]],
         grad_fn=<SoftmaxBackward>)
```

```
In [11]: probs, index = torch.max(sm_outputs, dim=1)

for p, i in zip(probs, index):
    print('{0} - {1:.4f}'.format(classes[i], p))
```

```
cat - 0.4080
ship - 0.3793
ship - 0.4816
ship - 0.7051
```


```
In [12]: total_correct = 0
total_images = 0
confusion_matrix = np.zeros([10,10], int)
with torch.no_grad():
    for data in testloader:
        images, labels = data
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total_images += labels.size(0)
        total_correct += (predicted == labels).sum().item()
        for i, l in enumerate(labels):
            confusion_matrix[l.item(), predicted[i].item()] += 1

model_accuracy = total_correct / total_images * 100
print('Model accuracy on {0} test images: {1:.2f}%'.format(total_images, model_accuracy))
```

```
Model accuracy on 10000 test images: 54.98%
```

```
In [14]: from matplotlib import pyplot as plt

fig, ax = plt.subplots(1,1,figsize=(8,6))
ax.matshow(confusion_matrix, aspect='auto', vmin=0, vmax=1000, cmap=plt.get_cm
ap('Blues'))
plt.ylabel('Actual Category')
plt.yticks(range(10), classes)
plt.xlabel('Predicted Category')
plt.xticks(range(10), classes)
plt.show()
```



```
In [19]: print('{0:10s} - {1}'.format('Category', 'Accuracy'))
for i, r in enumerate(confusion_matrix):
    print('{0:10s} - {1:.1f}'.format(classes[i], r[i]/np.sum(r)*100))
```

```
Category - Accuracy
plane - 44.1
car - 66.2
bird - 44.1
cat - 32.5
deer - 34.4
dog - 50.5
frog - 76.9
horse - 56.9
ship - 79.8
truck - 64.4
```

```
In [20]: print('actual/pred'.ljust(16), end='')
for i,c in enumerate(classes):
    print(c.ljust(10), end='')
print()
for i,r in enumerate(confusion_matrix):
    print(classes[i].ljust(16), end='')
    for idx, p in enumerate(r):
        print(str(p).ljust(10), end='')
    print()

    r = r/np.sum(r)
    print(''.ljust(16), end='')
    for idx, p in enumerate(r):
        print(str(p).ljust(10), end='')
    print()
```

actual/pred	plane	car	bird	cat	deer	dog	f
rog	horse	ship	truck				
plane	441	19	99	21	9	23	4
1	13	265	69				
		0.441	0.019	0.099	0.021	0.009	0.023
0.041	0.013	0.265	0.069				
car	18	662	6	17	3	3	2
4	6	64	197				
		0.018	0.662	0.006	0.017	0.003	0.003
0.024	0.006	0.064	0.197				
bird	46	11	441	66	73	100	1
60	30	52	21				
		0.046	0.011	0.441	0.066	0.073	0.1
0.16	0.03	0.052	0.021				
cat	10	10	79	325	48	214	2
09	29	42	34				
		0.01	0.01	0.079	0.325	0.048	0.214
0.209	0.029	0.042	0.034				
deer	21	3	197	41	344	86	1
72	83	47	6				
		0.021	0.003	0.197	0.041	0.344	0.086
0.172	0.083	0.047	0.006				
dog	12	1	85	174	30	505	1
13	40	27	13				
		0.012	0.001	0.085	0.174	0.03	0.505
0.113	0.04	0.027	0.013				
frog	3	8	59	35	36	36	7
69	16	14	24				
		0.003	0.008	0.059	0.035	0.036	0.036
0.769	0.016	0.014	0.024				
horse	12	4	51	70	54	120	6
7	569	13	40				
		0.012	0.004	0.051	0.07	0.054	0.12
0.067	0.569	0.013	0.04				
ship	58	31	31	18	2	6	1
6	4	798	36				
		0.058	0.031	0.031	0.018	0.002	0.006
0.016	0.004	0.798	0.036				
truck	16	110	17	30	2	8	5
3	11	109	644				
		0.016	0.11	0.017	0.03	0.002	0.008
0.053	0.011	0.109	0.644				


In [23]: net.conv1

Out[23]: Conv2d(3, 6, kernel\_size=(5, 5), stride=(1, 1))



```
In [29]: conv1 = net.conv1
weight = conv1.weight.data

fig, axes = plt.subplots(1, len(weight), figsize=(5,5))
for idx, image in enumerate(weight):
    axes[idx].imshow(convert_to_imshow_format(image))
    axes[idx].set_xticks([])
    axes[idx].set_yticks([])
```



```
In [52]: NUM_CLASSES = 10

class AlexNet(nn.Module):
    def __init__(self, num_classes=NUM_CLASSES):
        super(AlexNet, self).__init__()
        self.features = nn.Sequential(
            nn.Conv2d(3, 64, kernel_size=3, stride=2, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2),
            nn.Conv2d(64, 128, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2),
            nn.Conv2d(128, 256, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2),
            nn.Conv2d(256, 512, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2),
            nn.Conv2d(512, 1024, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2),
        )
        self.classifier = nn.Sequential(
            nn.Dropout(),
            nn.Linear(1024 * 2 * 2, 2048),
            nn.ReLU(inplace=True),
            nn.Dropout(),
            nn.Linear(2048, 2048),
            nn.ReLU(inplace=True),
            nn.Linear(2048, num_classes),
        )

    def forward(self, x):
        x = self.features(x)
        x = x.view(x.size(0), 1024 * 2 * 2)
        x = self.classifier(x)
        return x

alex_net = AlexNet()
```

```
In [53]: import torch.optim as optim

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(alex_net.parameters(), lr=0.001, momentum=0.9)
```

```
In [54]: import os

model_directory_path = '/home/CIFAR-10 Classifier Using CNN in PyTorch/model/'
model_path = model_directory_path + 'cifar-10-cnn-model_alex_net.pt'

if not os.path.exists(model_directory_path):
    os.makedirs(model_directory_path)

if os.path.isfile(model_path):
    # Load trained model parameters from disk
    alex_net.load_state_dict(torch.load(model_path))
    print('Loaded model parameters from disk.')
else:
    for epoch in range(2): # Loop over the dataset multiple times

        running_loss = 0.0
        for i, data in enumerate(trainloader, 0):
            # get the inputs
            inputs, labels = data

            # zero the parameter gradients
            optimizer.zero_grad()

            # forward + backward + optimize
            outputs = alex_net(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            # print statistics
            running_loss += loss.item()
            if i % 2000 == 1999: # print every 2000 mini-batches
                print('[%d, %5d] loss: %.3f' %
                      (epoch + 1, i + 1, running_loss / 2000))
                running_loss = 0.0
        print('Finished Training.')
        torch.save(alex_net.state_dict(), model_path)
        print('Saved model parameters to disk.')
```

```
[1, 2000] loss: 2.303
[1, 4000] loss: 2.261
[1, 6000] loss: 2.063
[1, 8000] loss: 1.906
[1, 10000] loss: 1.802
[1, 12000] loss: 1.728
[2, 2000] loss: 1.622
[2, 4000] loss: 1.568
[2, 6000] loss: 1.456
[2, 8000] loss: 1.418
[2, 10000] loss: 1.374
[2, 12000] loss: 1.324
Finished Training.
Saved model parameters to disk.
```

```
In [45]: dataiter = iter(testloader)
images, labels = dataiter.next()

fig, axes = plt.subplots(1, len(images), figsize=(12,2.5))
for idx, image in enumerate(images):
    axes[idx].imshow(convert_to_imshow_format(image))
    axes[idx].set_title(classes[labels[idx]])
    axes[idx].set_xticks([])
    axes[idx].set_yticks([])
```



```
In [46]: outputs = alex_net(images)
sm = nn.Softmax(dim=1)
sm_outputs = sm(outputs)
print(sm_outputs)

tensor([[0.0633, 0.0763, 0.1413, 0.1392, 0.0593, 0.1062, 0.1346, 0.0423, 0.13
77,
        0.1000],
        [0.1124, 0.3729, 0.0102, 0.0065, 0.0060, 0.0033, 0.0047, 0.0070, 0.24
56,
        0.2314],
        [0.1307, 0.1137, 0.1016, 0.0657, 0.0553, 0.0562, 0.1327, 0.0286, 0.19
48,
        0.1207],
        [0.2044, 0.0921, 0.0160, 0.0032, 0.0033, 0.0014, 0.0021, 0.0016, 0.64
81,
        0.0278]], grad_fn=<SoftmaxBackward>)
```

```
In [47]: probs, index = torch.max(sm_outputs, dim=1)

for p, i in zip(probs, index):
    print('{0} - {1:.4f}'.format(classes[i], p))

bird - 0.1413
car - 0.3729
ship - 0.1948
ship - 0.6481
```

```

In [48]: total_correct = 0
total_images = 0
confusion_matrix = np.zeros([10,10], int)
with torch.no_grad():
    for data in testloader:
        images, labels = data
        outputs = alex_net(images)
        _, predicted = torch.max(outputs.data, 1)
        total_images += labels.size(0)
        total_correct += (predicted == labels).sum().item()
    for i, l in enumerate(labels):
        confusion_matrix[l.item(), predicted[i].item()] += 1

model_accuracy = total_correct / total_images * 100
print('Model accuracy on {0} test images: {1:.2f}%'.format(total_images, model_accuracy))

```

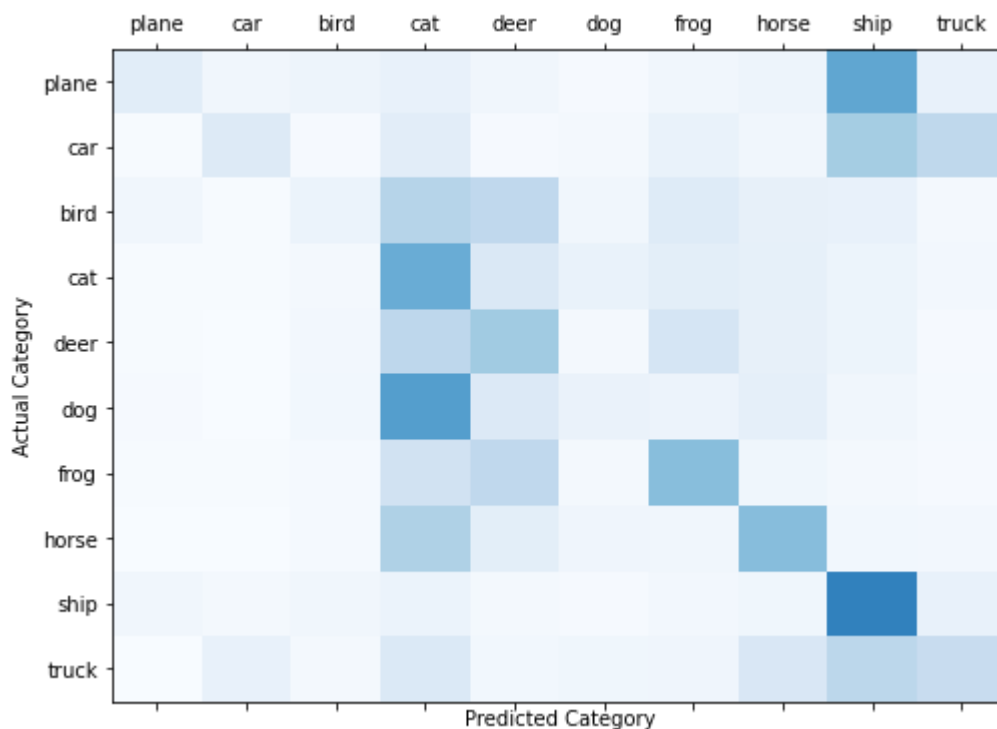
Model accuracy on 10000 test images: 30.14%

```

In [49]: from matplotlib import pyplot as plt

fig, ax = plt.subplots(1,1,figsize=(8,6))
ax.matshow(confusion_matrix, aspect='auto', vmin=0, vmax=1000, cmap=plt.get_cmap('Blues'))
plt.ylabel('Actual Category')
plt.yticks(range(10), classes)
plt.xlabel('Predicted Category')
plt.xticks(range(10), classes)
plt.show()

```



In [ ]: